

Pre-M1 (c)

*: dereferencer
&: referencer.

C / Memory

- `ptr = malloc(size);` Allocates uninitialized block, size size
(@ pointer `ptr` in memory)
- `ptr = calloc(num_items, size);` allocates num_items blocks of zeroed memory of size size
(@ pointer `ptr`)
- `free(ptr);` free memory allocated at `ptr`. do this if you'll never use that mem. again.
- `realloc(ptr, new_size);` Either
 - Expands or contracts memory allocated by `ptr` to `new_size`
 - Allocates a new block, copies the memory it can from old block, frees old block.

Alignment

- char: 1 byte, no alignment rules
- short: 2 bytes, aligned by halfword
- int: 4 bytes, aligned by word
- void*

Stack vs. Heap Space

- Stack space: memory allocated on the call stack
 - most primitive datatypes
- Heap space: memory allocated during execution
 - allocated memory
- Static storage: global vars, static literals
- Code storage: bytecode that comprises the code being run

Struct:

- multiple blocks of data in one structure.
- info is bunched adjacently & spaced according to alignment rules.

Unions:

- multiple data types stored in one block of memory.
- allocated to be the largest data type
- set Block to 0 = set everything to 0.

CALL: Compiler

Assembler

passes over twice
for PC-relative
labels.

translates C into a lower-level asm lang. like RISC-V

translates from Assembly into machine code.

↳ reads & user directives, which give directions to the assembler for how to handle mem, etc.

↳ object file format:

; header - size & pos of other pieces

; text segment - code

; data segment - binary rep of static data in source file

Linker

resolves all assembled symbols, possibly w/ libraries, etc.

and references to static data.

↳ stores this info in a relocation table & symbol table.

↳ lowest part of CALL

Loader

Running the executable, putting data into memory.

Is the OS

Control Logic

- PCSel: If ≠ branch or jump, PC = label, not PC+4.

- ImmSel: diff instruction types have diff ways of making immediates
 - ↳ immSel determines which to use.

- RegWEN: do we write to register? No for store & branch

- BSEL, ASel: picking what to add to ALU. ALU PC & RS1, while B bwn imm & RS2

- ALU Sel: which ALU op

- MemRW: write to memory?

- WB Sel: What to write to dest reg? Memory, ALU, or PC+4?

Number Rep:

Unsigned

- representable range: $[0, 2^n - 1]$
- just ignores signs, best for nonnegative-only metrics like length.

- if we wanted signs..,

Signed magnitude

- Representable Range: $[-2^{n-1}, 2^{n-1} - 1]$
- one bit determining sign, rest comprising int's as normal

bias

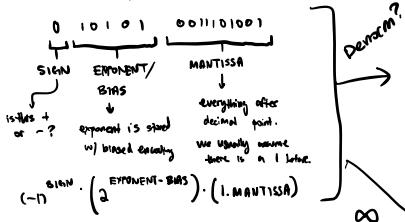
- Representable Range: $[B, 2^n - 1 + B]$, usually $B > 0$.
- allows us to represent neg. #'s as nonneg.
- would need to perform bias externally and it would work like unsigned.

2's complement

- Representable Range: $[2^{n-1}, 2^{n-1} - 1]$
- then, to repr. a negative number, flip all bits in original number & add 1 to the result.

Floating Point

Four Components



Change behavior of 0.b00...0 st. range extends to exactly 0.
↳ instead of this being 1.2 bits, make it $0.2^{B_{\text{bias}}+1}$ when exponent = 0
↳ decreases precision but at least we can represent #'s < 1

If exponents are all 1's, then our value is ∞
only if mantissa = 0. Otherwise, is NaN.

MTTF / MTTR

mean time to failure: higher = better
mean time to repair: lower the better

$$MTBF = MTTF + MTTR, \text{ Availability} = \frac{MTTF}{MTBF}$$

if we have smth running for X days out of Y, avail. = X/Y . Use this to find MTTF / MTTR.
If we have one, can find the other.

RAID

- RAID 0: just put all memory inside one
 - not dependable

- RAID 1: copy disk n times
 - restricts amt. of storage we can have

- RAID 4: store info across disks w/ parity bit.
 - can run ECC, and more mem efficient.

- RAID 5: same as 4, but stratified.

ECC

→ parity = even-odd of sum of the bits, found via XOR.

→ Hamming Distance = # bits different b/w two strings.

→ Hamming ECC: Use library value to determine corrupted bit.

↳ Check if parity bits need to be 1 or 0 after adding infill bits & encoding.

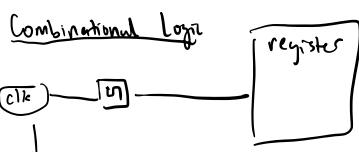
↳ if p's have errors, find bit overlap errors & flip the sf.



Pre-Midterm (ASM)

Syntax:

- upper immediate is anything $\geq 2^8 = 1000$.
 - this is basically adding values above 12th bit to an index.
 - example adds upper immediate of an offset to the PC.
- jalr stores pc+4 in a register. commonly use "ra" to find where we jump back
- calling conventions!!
 - values of a, t registers do not carry over.
 - values of s registers & ra must be stored on stack, done w/ sp.



clock repeatedly turns on & off. this is the clock cycle.
 Registers only update once per clock cycle, when clk turns on. however, you must wait until after clk-to-q timer to see updates.

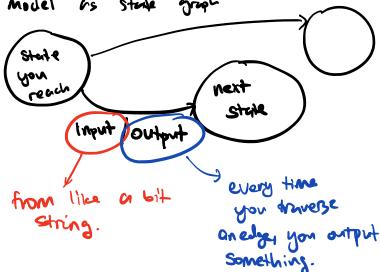
- each component in the circuit has a propagation time.
- setup time: amount of time before a clock interval that input must be stable (constant)
- hold time: amount of time after a clock interval that input must be stable. depends on components.
- $t_{clk-to-q} + t_{shortest-combo-path} \leq t_{hold}$
- $t_{clk-to-q} + t_{longest-combo-path} + t_{setup} \leq t_{clock-period}$.

SHORTEST comb. path is determined by finding least time between two timed components. (NOT necessarily reg's!)

LONGEST comb. path is determined by finding the longest time between two timed components

Finite States

- detect a bit sequence or state
- model as state graph



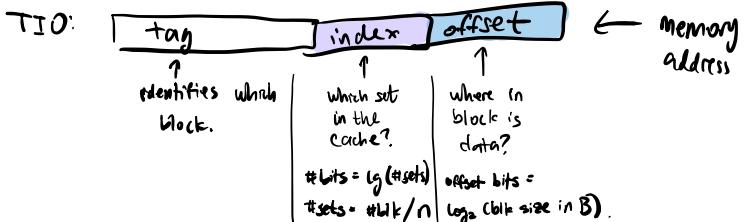
Caches

- intermediate b/w energy cost & time for registers & DRAM.

$$\lg(\text{mem size}) = \text{addr. bit width} = |T| + |I| + |O|$$

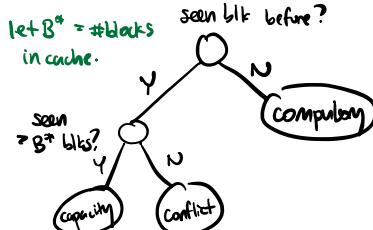
$$\text{Cache size} = (\text{blk size})(\# \text{blk}) = N \cdot 2^{101} + |I|$$

- n-way set associativity: we can put n addresses in a certain index
 - fully associative: there is no index, we just put addresses in the cache.
 - more associativity: ↓ hit speed, ↑ hit time, ↑ hit rate.
 - however, cache misses are much more costly than cache hits, so a respectable hit rate is crucial.
 - direct-mapped = one block per index.



Cache Misses

- 1) compulsory miss — have not yet fetched block.
- 2) capacity miss — req'd data evicted because no space left.
- 3) conflict miss — need to update cache to fit the tag.



Eviction Policy

- FIFO: first block added is evicted
- LRU: least recently used block is evicted
 - b/c of locality, this is better.

AMAT = avg. memory access time.

$$(\text{hit time}) + (\text{miss rate}) \cdot (\text{miss penalty})$$

Multi-level Caches

↳ usually increasing in associativity.

L1 cache → L2 cache → etc.

$$\text{AMAT} = \text{L1 HT} + \text{L1 MR} \cdot (\text{L2 HT} + \text{L2 MR} \cdot (\text{L3 ...}))$$

Post - MT:

Pipeline: (Data path in ref sheet)

- latency: delay to process operation
- throughput: # ops done in one op.
- pipelining increases latency, but also throughput.

There are 5 stages:

- IF \rightarrow instruction fetch. what's the instruction
- ID \rightarrow instruction decode - find the meaning of instruction
- EX \rightarrow execute - do stuff w/ branch comp, ALU, etc.
- MEM \rightarrow memory - put info into memory
- WB \rightarrow write back - what to reinsert into memory?

Must put registers between stages

Hazards:

- structural hazard
 - hardware doesn't support running these insts. @ same time,
 - solutions are either stalling or + HW, and cur. regfile implementation seems to work.
- data hazard
 - instructions have data dependency,
 - need to wait for memory to behave as intended before using that block.
 - 1) register access — same reg. written & read in 1 cycle.
solution: write, then read
 - 2) ALU result — inst. depends on WB's register write from previous.
solution: forwarding: grab info from pipeline stage, requiring add'l hardware.
 - 3) load data hazard — needs stall.
need to access memory only changed during end of clock cycle, @ beginning of this cycle.
 \hookrightarrow need load delay, perf. can be salvaged by code scheduling.

Control Hazard

- how to update the pipeline after branch is taken
 - \hookrightarrow must kill all processes and start @ label if branch taken.
 - Sol. to speedup is to take a branch predictor.
 - Loss of 0 cc if correct, 3CC if wrong.

Virtual Memory

Purpose: illusion of isolated processes, important for security.

- \hookrightarrow Memory is stored at fixed-size pages.
- \hookrightarrow can spoof addr. of physical memory w/ virtual memory.
- \hookrightarrow can have arbitrary amounts of physical & virtual memory.
 - \hookrightarrow each index can be represented as [P/N]PN \rightarrow offset.
- \hookrightarrow memory manager translates, i.e.

TLB: translation lookaside buffer.

- \hookrightarrow cache for VPN-PPN mappings.

Parallelism: Can use the fact that CPU has multiple cores to do multiple things @ once

SIMD: does vector addition, in order to save time for array addition. should have a speed increase proportional to vector size.

from Intel AVX library.

data-level parallelism

Amdahl's Law: Speedup bounded by anti. not ||izable.

$$\frac{1}{(1-F) + \frac{F}{S}}, \text{ where } F \text{ is \% code sped up. } S \text{ = sped factor.}$$

Open MP:

- split a looped procedure into a series of threads:

- #pragma omp parallel
 - \hookrightarrow runs n threads in following block
- #pragma omp for
 - \hookrightarrow splits for loop automatically
- #pragma critical
 - \hookrightarrow prevents data races.

Multiple threads receiving + modifying same data @ unsync'd times

Whenever there's C code:
for example, $x = x + 1$
this is a load, add, and store. All 3 actions can happen at different times, with any # of others in between.

- #pragma omp reduction \rightarrow compute the aggregate of some value in parallel

- run multiple processes = multiple independent instances.
- worker & manager processes.

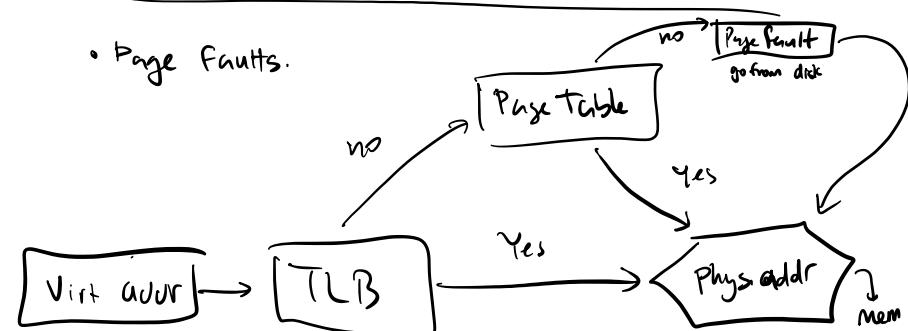
Manager: assign work & inform user of progress
Worker: performs task from manager.

Needs:

- Code that sends tasks w/ procID = D.
- Code that recvs tasks, every other process.

\hookrightarrow Want to minimize size of message assigning new work.

- Page faults.



OpenMPI

MAP REDUCE & SPARK

- Many workers do small parts of large tasks.

map(*in-key*, *in-value*)

↳ (*int-key*, *int-value*)

reduce(*int-key*, *int-value*)

↳ (*out-value*)

↳ continuation of multiprocessing section above.

• initialize multiprocessing: MPI-Init(*argc*, *argv*)

• get # of nodes / proc ID (*MPI_Comm_size(comm, &size)*; *MPI_Comm_rank(comm, &rank)*)

• functions taken in pointers & values written to pointers;

• returns # of processes with successful.

• can use *MPI_COMM_WORLD* for common value.

• run acc. to worker / manager model?

• MPI-Recv(*buf*, *count*, *datatype*, *source*, *tag*, *comm*, *status*)

• MPI-Send(*buf*, *count*, *datatype*, *dest*, *tag*, *comm*)

• *buf*: array of data to be sent / buffer to receive data

• *datatype*: constants used to specify input type

• *count*: # elements of datatype to receive

• *dest*: forward, the procID of intended msg recipient

• *source*: for recv, the procID of expected msg sender.

• clean up w/ MPI-Finalize()

• *tag*: for when you want to further classify msgs.
o send/recv only matches when tags identical

• *comm*: communication group; for this class
set w/ MPI-Comm-World

• *status*: for recv, contains the source/tag, error

msg. of communication.

HOW MUCH faster?

reg/func
inlining: < 2x



loop
unrolling: < 2x



cache opt: ~10x



SIMD: ~8x



openMP: #cores/node



openMPI: #cores.



Pros?

• easy change

• reduces # accesses

• reduce branching

• surprisingly good

• fairly applicable
• minimal overhead

• more flexible than SIMD
or MPI

• can be extended
arbitrarily large

Cons?

• minimal effect

• compiler might not do

• minimal effect

• penalty to maintainability

• req's alg changes

• Amdahl
• restricted by HW.

• concurrency issues
• high overhead

• expensive comm.
• high overhead.