

Lecture 4

Kernel Mode Transfers

- Syscall
 - Problem requests a system service, such as exit
 - Like a function call, but **outside** the process
 - Doesn't have the address of the system function to call
 - Like a Remote Procedure Call (RPC) -- arguments must then be passed in
 - Triggered by the user process
 - Marshall the syscall id and args in registers and exec syscall
- Interrupt
 - External async event triggers context switch
 - e.g. timer, i/o device.
 - Independent of user process
- Trap/Exception
 - Synchronous event in process triggers context switch
 - Most likely some error

Interrupts

- Interrupt processing not visible to the user process
 - Occurs between instructions, restarted transparently
- Interrupt controller / masking
 - Interrupts invoked by interrupt lines on devices
 - Controller chooses which interrupt requests to honor
 - The mask enables/disables the interrupt
 - The priority encoder picks the highest enabled interrupt
 - Software Interrupt will be set and cleared by the software
 - CPU can disable all interrupts w/ an internal flag
 - Non-Maskable Interrupt (NMI) line **can't** be disabled.

- Interrupt vector
 - Limited number of entry points into kernel
- Kernel interrupt stack
 - Cannot share stacks with the user.
 - By separating the stack, handler works regardless of state of user code
- Atomic transfer of control
 - "Single instruction" - like to change:
 - program counter
 - stack pointer
 - mem protection
 - kernel/user mode

During a syscall / interrupt, we will save stack from user memory onto the kernel exception stack

Managing Processes

- How to manage state of process?
 - Create
 - Exit
- Everything outside of the kernel is running in a process
- Processes are created and managed by processes

Bootstrapping: First process (init, pid 0) is started by the kernel, after which all processes will be created by other processes

API:

- `exit` - terminate a process (`pthread_exit`)
- `fork` - copy current process (`pthread_create`)
- `exec` - change program being run by current proc
- `wait` - wait for proc to finish (`pthread_join`)
- `kill` - send a signal (interrupt-like notif) to another proc

- `sigaction` - set handlers for signals

Forking

- `pid_t fork()` -- copy the current process
 - new process has different pid
 - new process contains one thread
- return value from `fork()` : `pid` (int)
 - when > 0 :
 - running in original (parent) process
 - return value is `pid` of new child
 - when $= 0$:
 - running in new child process
 - when < 0 :
 - error, must be handled somehow
 - running in original process
- fork race
 - prints things in a random order
 - sleeping each process doesn't help

Sigaction

Needs a signal handler

Common POSIX signals:

- `SIGINT` : ctrl-C
- `SIGTERM` : `kill` shell command
- `SIGSTOP` : ctrl-Z
- `SIGKILL` , `SIGSTOP` : terminate/stop program. can't be changed with `sigaction`

Files

Unix/POSIX idea: **everything** is a file. There are identical interfaces for:

- files on disk

- devices (terminals, printers, etc)
- regular files on disk
- networking sockets

Based on the system calls `open` , `read` , `write` , and `close` . Additionally, `ioctl()` for custom configuration that doesn't quite fit the above.

- was radical when proposed, but hasn't been changed since.
- so they prolly did something right.

Every process has a current working directory (cwd). Absolute paths ignore it, and relative paths are relative to it. `.` is the cwd itself, and `..` is the parent dir.

- can be set with a syscall

File API: Operates on "streams" -- unformatted sequences of bytes with a position.

- Open streams are represented by a pointer to a `FILE` ds
- Error is reported with a null response
- Three predefined streams, `stdin` , `stdout` , and `stderr` , are opened implicitly when prog is executed
- All three can be redirected